

In [1]:

```

1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6 import warnings
7 warnings.filterwarnings('ignore')

```

In [2]:

```

1 df1=pd.read_csv('test.csv')
2 df2=pd.read_csv('train.csv')

```

In [3]:

```
1 df=df1.append(df2)
```

In [4]:

```
1 df.head()
```

Out[4]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coap
0	LP001015	Male	Yes	0	Graduate	No	5720	
1	LP001022	Male	Yes	1	Graduate	No	3076	
2	LP001031	Male	Yes	2	Graduate	No	5000	
3	LP001035	Male	Yes	2	Graduate	No	2340	
4	LP001051	Male	No	0	Not Graduate	No	3276	

In [5]:

```
1 df.shape
```

Out[5]:

(981, 13)

In [6]:

```
1 df.describe()
```

Out[6]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	981.000000	981.000000	954.000000	961.000000	902.000000
mean	5179.795107	1601.916330	142.511530	342.201873	0.835920
std	5695.104533	2718.772806	77.421743	65.100602	0.370553
min	0.000000	0.000000	9.000000	6.000000	0.000000
25%	2875.000000	0.000000	100.000000	360.000000	1.000000
50%	3800.000000	1110.000000	126.000000	360.000000	1.000000
75%	5516.000000	2365.000000	162.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

In [7]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               981 non-null    object
1   Gender                957 non-null    object
2   Married               978 non-null    object
3   Dependents            956 non-null    object
4   Education             981 non-null    object
5   Self_Employed         926 non-null    object
6   ApplicantIncome       981 non-null    int64
7   CoapplicantIncome     981 non-null    float64
8   LoanAmount            954 non-null    float64
9   Loan_Amount_Term      961 non-null    float64
10  Credit_History         902 non-null    float64
11  Property_Area         981 non-null    object
12  Loan_Status           614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 107.3+ KB
```

In [8]:

```
1 df.isnull().sum()
```

Out[8]:

```
Loan_ID      0
Gender       24
Married       3
Dependents   25
Education     0
Self_Employed 55
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   27
Loan_Amount_Term 20
Credit_History 79
Property_Area 0
Loan_Status  367
dtype: int64
```

In [9]:

```
1 df.isnull().sum()/len(df)*100
```

Out[9]:

```
Loan_ID      0.000000
Gender       2.446483
Married       0.305810
Dependents   2.548420
Education     0.000000
Self_Employed 5.606524
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount   2.752294
Loan_Amount_Term 2.038736
Credit_History 8.053007
Property_Area 0.000000
Loan_Status  37.410805
dtype: float64
```

In [10]:

```
1 mode_sf_emp=df['Self_Employed'].mode()
```

In [11]:

```
1 mode_sf_emp
```

Out[11]:

```
0    No
Name: Self_Employed, dtype: object
```

In [12]:

```
1 df['Self_Employed']=df['Self_Employed'].fillna(method='ffill')
```

In [13]:

```
1 mode_ln_st=df['Loan_Status'].mode()  
2 mode_ln_st
```

Out[13]:

```
0    Y  
Name: Loan_Status, dtype: object
```

In [14]:

```
1 df['Loan_Status']=df['Loan_Status'].fillna(method='ffill')
```

In [15]:

```
1 median_cr=df['Credit_History'].median()  
2 median_cr
```

Out[15]:

```
1.0
```

In [16]:

```
1 df['Credit_History']=df['Credit_History'].fillna(median_cr)
```

In [17]:

```
1 df.isnull().sum()/len(df)*100
```

Out[17]:

```
Loan_ID          0.000000  
Gender           2.446483  
Married          0.305810  
Dependents       2.548420  
Education        0.000000  
Self_Employed    0.000000  
ApplicantIncome  0.000000  
CoapplicantIncome 0.000000  
LoanAmount       2.752294  
Loan_Amount_Term 2.038736  
Credit_History   0.000000  
Property_Area     0.000000  
Loan_Status      37.410805  
dtype: float64
```

In [18]:

```
1 df.dropna(inplace=True)
```

In [19]:

```
1 df.isnull().sum()
```

Out[19]:

```
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed    0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History  0
Property_Area    0
Loan_Status      0
dtype: int64
```

In [20]:

```
1 #Numerical_col
2 num_col=df.select_dtypes(include=['int64','float64']).columns
3 num_col
```

Out[20]:

```
Index(['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
      'Loan_Amount_Term', 'Credit_History'],
      dtype='object')
```

In [21]:

```
1 #Categorical_col
2 cat_col=df.select_dtypes(include=['O']).columns
3 cat_col
```

Out[21]:

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
      'Self_Employed', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

In [22]:

```
1 df.drop('Dependents',axis=1,inplace=True)
```

In [23]:

```
1 df.head(10)
```

Out[23]:


	Loan_ID	Gender	Married	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
1	LP001003	Male	Yes	Graduate	No	4583	1508.0
2	LP001005	Male	Yes	Graduate	Yes	3000	0.0
3	LP001006	Male	Yes	Not Graduate	No	2583	2358.0
4	LP001008	Male	No	Graduate	No	6000	0.0
5	LP001011	Male	Yes	Graduate	Yes	5417	4196.0
6	LP001013	Male	Yes	Not Graduate	No	2333	1516.0
7	LP001014	Male	Yes	Graduate	No	3036	2504.0
8	LP001018	Male	Yes	Graduate	No	4006	1526.0
9	LP001020	Male	Yes	Graduate	No	12841	10968.0
10	LP001024	Male	Yes	Graduate	No	3200	700.0

In [24]:

```
1 sns.countplot(df['Loan_Status'])
```

Out[24]:

<AxesSubplot:xlabel='Loan_Status', ylabel='count'>




In [25]:

```
1 sns.countplot(df['Gender'],hue=df['Loan_Status'])
```

Out[25]:

<AxesSubplot:xlabel='Gender', ylabel='count'>




In [26]:

```
1 sns.countplot(df['Self_Employed'],hue=df['Loan_Status'])
```

Out[26]:

<AxesSubplot:xlabel='Self_Employed', ylabel='count'>




In [27]:

```
1 sns.countplot(df['Property_Area'],hue=df['Loan_Status'])
```

Out[27]:

<AxesSubplot:xlabel='Property_Area', ylabel='count'>




In [28]:

```
1 sns.scatterplot(df['ApplicantIncome'],df['LoanAmount'])
```

Out[28]:

<AxesSubplot:xlabel='ApplicantIncome', ylabel='LoanAmount'>

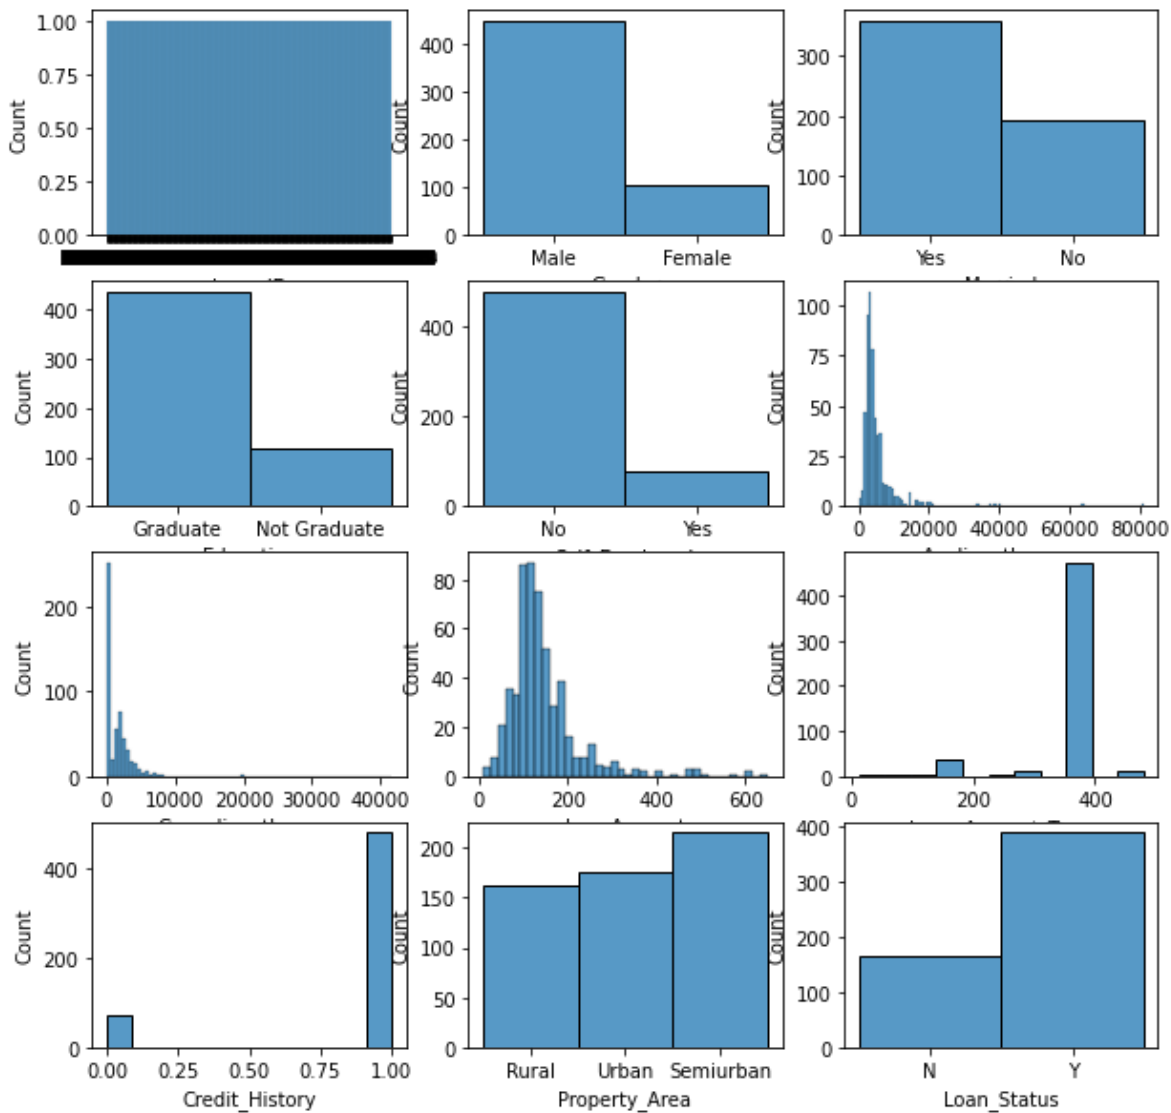


In [29]:

```

1 plt.figure(figsize=(10,10))
2 count=1
3 for i in df:
4     plt.subplot(4,3,count)
5     sns.histplot(x=df[i],data=df)
6     count+=1

```




In [30]:

```
1 sns.distplot(df['ApplicantIncome'])
```

Out[30]:

<AxesSubplot:xlabel='ApplicantIncome', ylabel='Density'>

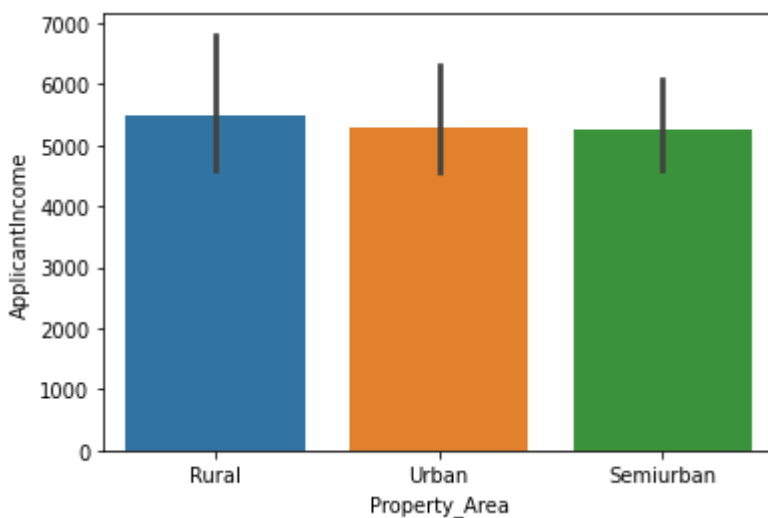


In [31]:

```
1 sns.barplot(df['Property_Area'],df['ApplicantIncome'])
```

Out[31]:

<AxesSubplot:xlabel='Property_Area', ylabel='ApplicantIncome'>




In [32]:

```
1 sns.distplot(df['LoanAmount'])
```

Out[32]:

<AxesSubplot:xlabel='LoanAmount', ylabel='Density'>




In [33]:

```
1 plt.figure(figsize=(15,11))
2 sns.heatmap(df.corr(),annot=True)
```

Out[33]:

<AxesSubplot:>



LabelEncoder

In [34]:

```
1 from sklearn.preprocessing import LabelEncoder
```

In [35]:

```
1 le=LabelEncoder()
```

In [36]:

```
1 cat_col=df.select_dtypes(include='O').columns
2 cat_col
```

Out[36]:

```
Index(['Loan_ID', 'Gender', 'Married', 'Education', 'Self_Employed',
      'Property_Area', 'Loan_Status'],
      dtype='object')
```

In [37]:

```
1 for i in cat_col:
2     df[i]=le.fit_transform(df[i])
```

In [38]:

```
1 df.head()
```

Out[38]:

Loan_ID	Gender	Married	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
1	0	1	1	0	4583	1508.0
2	1	1	1	0	3000	0.0
3	2	1	1	1	2583	2358.0
4	3	1	0	0	6000	0.0
5	4	1	1	0	5417	4196.0

In [39]:

```
1 x=df.drop('Loan_Status',axis=1)
2 x
```

Out[39]:

Loan_ID	Gender	Married	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
1	0	1	1	0	4583	1508.0
2	1	1	1	0	3000	0.0
3	2	1	1	1	2583	2358.0
4	3	1	0	0	6000	0.0
5	4	1	1	0	5417	4196.0
...
609	548	0	0	0	2900	0.0
610	549	1	1	0	4106	0.0
611	550	1	1	0	8072	240.0
612	551	1	1	0	7583	0.0
613	552	0	0	0	4583	0.0

553 rows × 11 columns

In [40]:

```
1 y=df['Loan_Status']  
2 y
```

Out[40]:

```
1      0  
2      1  
3      1  
4      1  
5      1
```

..

```
609    1  
610    1  
611    1  
612    1  
613    0
```

Name: Loan_Status, Length: 553, dtype: int32

train_test_split

In [41]:

```
1 from sklearn.model_selection import train_test_split
```

In [42]:

```
1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=1)
```

In [43]:

```
1 X_train.shape
```

Out[43]:

(442, 11)

In [44]:

```
1 X_test.shape
```

Out[44]:

(111, 11)

In [45]:

```
1 y_train.shape
```

Out[45]:

(442,)

In [46]:

```
1 y_test.shape
```

Out[46]:

(111,)

LogisticRegression

In [47]:

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import classification_report
```

In [48]:

```
1 def my_model(clf):
2     clf.fit(X_train,y_train)
3     y_train_pred=clf.predict(X_train)
4     y_test_pred=clf.predict(X_test)
5     print('Train Data')
6     print(classification_report(y_train,y_train_pred))
7     print('Test Data')
8     print(classification_report(y_test,y_test_pred))
```

In [49]:

```
1 lr=LogisticRegression()
```

In [50]:

```
1 my_model(lr)
```

Train Data

	precision	recall	f1-score	support
0	0.85	0.38	0.53	133
1	0.79	0.97	0.87	309
accuracy			0.79	442
macro avg	0.82	0.68	0.70	442
weighted avg	0.80	0.79	0.77	442

Test Data

	precision	recall	f1-score	support
0	0.62	0.42	0.50	31
1	0.80	0.90	0.85	80
accuracy			0.77	111
macro avg	0.71	0.66	0.67	111
weighted avg	0.75	0.77	0.75	111

DecisionTreeClassifier

In [51]:

```
1 from sklearn.tree import DecisionTreeClassifier
```

In [52]:

```
1 dt=DecisionTreeClassifier()
```

In [53]:

```
1 dt
```

Out[53]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

In [54]:

```
1 my_model(dt)
```

Train Data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	133
1	1.00	1.00	1.00	309
accuracy			1.00	442
macro avg	1.00	1.00	1.00	442
weighted avg	1.00	1.00	1.00	442

Test Data

	precision	recall	f1-score	support
0	0.49	0.55	0.52	31
1	0.82	0.78	0.79	80
accuracy			0.71	111
macro avg	0.65	0.66	0.66	111
weighted avg	0.72	0.71	0.72	111

In [55]:

```
1 from sklearn.model_selection import RandomizedSearchCV
```

In [56]:

```
1 param_grid={
2     'criterion':['gini','entropy'],
3     'class_weight':[None,'balanced'],
4     'max_depth':np.arange(2,50),
5     'min_samples_split':np.arange(2,50,2),
6     'min_samples_leaf':np.arange(2,50)
7 }
```



In [57]:

```
1 dt_rcv=RandomizedSearchCV(dt,param_distributions=param_grid,n_iter=10,scoring='f1',n_jo
```

In [58]:

```
1 dt_rcv.fit(X_train,y_train)
```

Out[58]:



In [59]:

```
1 dt_rcv.best_params_
```

Out[59]:

```
{'min_samples_split': 12,
 'min_samples_leaf': 41,
 'max_depth': 29,
 'criterion': 'entropy',
 'class_weight': None}
```

In [60]:

```
1 dt1=DecisionTreeClassifier(criterion='gini',class_weight='balanced',max_depth=39,min_sa
```

In [61]:

```
1 my_model(dt1)
```

Train Data					
	precision	recall	f1-score	support	
	0	0.65	0.89	0.75	133
	1	0.94	0.79	0.86	309
	accuracy			0.82	442
	macro avg	0.80	0.84	0.81	442
	weighted avg	0.85	0.82	0.83	442
Test Data					
	precision	recall	f1-score	support	
	0	0.42	0.71	0.53	31
	1	0.85	0.62	0.72	80
	accuracy			0.65	111
	macro avg	0.64	0.67	0.62	111
	weighted avg	0.73	0.65	0.67	111

RandomForestClassifier

In [62]:

```
1 from sklearn.ensemble import RandomForestClassifier
```

In [63]:

```
1 rf=RandomForestClassifier()
```

In [64]:

```
1 my_model(rf)
```

Train Data					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	133	
1	1.00	1.00	1.00	309	
accuracy			1.00	442	
macro avg	1.00	1.00	1.00	442	
weighted avg	1.00	1.00	1.00	442	
Test Data					
	precision	recall	f1-score	support	
0	0.71	0.55	0.62	31	
1	0.84	0.91	0.87	80	
accuracy			0.81	111	
macro avg	0.77	0.73	0.75	111	
weighted avg	0.80	0.81	0.80	111	

In [65]:

```
1 rf_rcv=RandomizedSearchCV(rf,param_distributions=param_grid,n_iter=10,scoring='f1',n_jobs=4)
```

In [66]:

```
1 rf_rcv.fit(X_train,y_train)
```

Out[66]:

RandomizedSearchCV

estimator: RandomForestClassifier

RandomForestClassifier

In [67]:

```
1 rf_rcv.best_params_
```

Out[67]:

```
{'min_samples_split': 38,
 'min_samples_leaf': 3,
 'max_depth': 27,
 'criterion': 'gini',
 'class_weight': None}
```

In [68]:

```
1 rf1=RandomForestClassifier(criterion='gini',class_weight='balanced',max_depth=31,min_s
```

In [69]:

```
1 my_model(rf1)
```

Train Data

	precision	recall	f1-score	support
0	0.62	0.69	0.65	133
1	0.86	0.82	0.84	309
accuracy			0.78	442
macro avg	0.74	0.75	0.74	442
weighted avg	0.79	0.78	0.78	442

Test Data

	precision	recall	f1-score	support
0	0.45	0.68	0.54	31
1	0.84	0.68	0.75	80
accuracy			0.68	111
macro avg	0.65	0.68	0.64	111
weighted avg	0.73	0.68	0.69	111

AdaBoostClassifier

In [70]:

```
1 from sklearn.ensemble import AdaBoostClassifier
```

In [71]:

```
1 adb=AdaBoostClassifier(n_estimators=450)
```

In [72]:

```
1 my_model(adb)
```

Train Data

	precision	recall	f1-score	support
0	0.95	0.80	0.87	133
1	0.92	0.98	0.95	309
accuracy			0.93	442
macro avg	0.93	0.89	0.91	442
weighted avg	0.93	0.93	0.93	442

Test Data

	precision	recall	f1-score	support
0	0.62	0.65	0.63	31
1	0.86	0.85	0.86	80
accuracy			0.79	111
macro avg	0.74	0.75	0.75	111
weighted avg	0.79	0.79	0.79	111

In [73]:

```
1 param_grid_ada={
2     'learning_rate':[0.1,0.01,1,2,3],
3     'n_estimators':[50,100,150]
4 }
```


In [76]:

```
1 adb_rcv=RandomizedSearchCV(adb,param_distributions=param_grid_ada,n_iter=10,scoring='f1')
```

In [77]:

```
1 adb_rcv.fit(X_train,y_train)
```

Out[77]:



In [78]:

```
1 adb_rcv.best_params_
```

Out[78]:

```
{'n_estimators': 50, 'learning_rate': 0.01}
```

In [81]:

```
1 adb1=AdaBoostClassifier(n_estimators=50,learning_rate=1)
```

In [82]:

```
1 my_model(adb1)
```

Train Data					
	precision	recall	f1-score	support	
0	0.89	0.57	0.70	133	
1	0.84	0.97	0.90	309	
accuracy			0.85	442	
macro avg	0.87	0.77	0.80	442	
weighted avg	0.86	0.85	0.84	442	
Test Data					
	precision	recall	f1-score	support	
0	0.70	0.61	0.66	31	
1	0.86	0.90	0.88	80	
accuracy			0.82	111	
macro avg	0.78	0.76	0.77	111	
weighted avg	0.81	0.82	0.82	111	

Best Prediction is given by AdaBoost Classifier model. 🏆

In []:

```
1
```